

Cost-Efficient VM Configuration Algorithm in the Cloud using Mix Scaling Strategy

Presenter: Li Lu

Li Lu, Jiadi Yu, Yanmin Zhu, Guangtao Xue, Shiyu Qian, Minglu Li

Department of Computer Science and Engineering,

Shanghai Jiao Tong University



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Popularity of Cloud Computing



VS.



- Cloud Computing vs. Typical Infrastructure
- Thanks to pay-per-use pricing, more elastic in management
 - Cloud computing can satisfy the peak workload without over-provision computing resources
 - e.g., Brickfish migrates its services to cloud leading to a decrease of cost from \$700,000 to \$200,000



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Difficulties in Managing Cloud Resources

➤ VM instance type selection

- Different VM instance type configurations → different performance & cost

➤ Precise VM instance type selection

- need accurate prediction of future workload (**difficult!**)
- even experienced administrators cannot precisely select VM instance type

➤ Key point: the **tradeoff** between **cost** and **performance** during the runtime

Region: US West (N. California) ▼

Operating system: ☒ Windows ☐ Linux ☐ My Images

Image: Microsoft Windows Server 2012 R2 Base (ami-cfa5b68a) ▼


Family: Compute optimized ▼ ☐ Show previous generations

Instance type: c4.xlarge ▼ vCPUs: 4 Memory: 7.5 GB



Existing Solutions

- Cost-aware **homogeneous** VM configurations
 - Same VM instance type
- **Multi-mechanisms** in VM configurations
 - Local-resize, replication, migration
- However, during the runtime in cloud,
 - Utilizing heterogeneous VM instance types is more cost-efficient
 - Migration of VM leads to high performance degradation



Outline

- Problem Definition
- Cost-efficient Mix Scaling Algorithm
- Evaluation
- Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



VM Configuration Model

- objective: minimize the renting cost of cloud resources
- constraints: the service rate of the configuration should be larger than the arrival rate of requests

$$\begin{aligned} \min \quad & \sum_{i=1}^K x_i c_i \\ \text{s.t.} \quad & \sum_{i=1}^K x_i \mu_i \geq \lambda \\ & x_i \in N, \quad i = 1, 2, \dots, K \end{aligned}$$

- the number of VM instance types: K
- the cost of the i^{th} VM instance type: c_i
- the maximum service rate of i^{th} VM instance type: μ_i
- the arrival rate of requests: λ
- the number of i^{th} VM instance type in the configuration: x_i



Differences between Two Constitute Configurations

- Due to the workload fluctuation, the two constitute VM configurations x_{old} and x_{new} are almost always different in all time slots.
 - Note that x_{old} and x_{new} are K-dimension vectors
- 3 situations may occur:
 - $x_{new} \geq x_{old}$: **more** VMs of **all types** are needed to meet performance requirement
 - $x_{new} \leq x_{old}$: **less** VMs of **all types** are needed to be cost-efficient
 - $x_{new} \neq x_{old}$: need to add or delete several VMs of **different instance types**
- For the first 2 situations, renting more or deleting several VMs would be OK
- For the 3rd situation, migrations would occur, which should be control to improve the performance

Cost-Migration Delay Tradeoff

➤ Tradeoff: **Cost** vs. **Migration delay**

- For **Cost**: the objective minimizes the cost

$$\min \sum_{i=1}^K x_i c_i$$

- For **Migration delay**: need to modeled



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY





Migration Delay Modeling

- Migration Mechanism in Cloud
 - Instead of directly migration, migration in cloud should **utilize the image server as a bridge**

- Migration Delay can be modelled as:

$$\alpha = 2 \frac{D}{b} + s$$

- where D is the image size, b is the bandwidth, s is the start time of a new VM



Cost-Migration Delay Tradeoff (COMDT) Problem

$$\begin{aligned} \min \quad & \sum_{i=1}^K x_i c_i \\ \text{s.t.} \quad & \sum_{i=1}^K x_i \mu_i \geq \lambda \\ & x_i \in N, \quad i = 1, 2, \dots, K \end{aligned}$$

Original Problem



$$\begin{aligned} \min \quad & \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \sum_{i=1}^K x_i(t) c_i \\ \text{s.t.} \quad & \sum_{i=1}^K x_i(t) \mu_i \geq \lambda(t), \forall t \\ & \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \alpha(t) \leq MT \\ & x_i \in N, \quad \forall i, t \end{aligned}$$

Migration Delay
Constraint

Cost-Migration Delay
Tradeoff Problem



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Outline

- Problem Definition
- Cost-efficient Mix Scaling Algorithm
- Evaluation
- Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Difficulty in Solving the COMDT Problem

- The COMDT problem aims to
 - minimize the long-term cost
 - constrain the long-term migration delay
- Notice that there are **two limits** in the objective and the migration delay constraint
 - Hard to solve with typical optimization techniques
 - Adopt **Lyapunov optimization** techniques

$$\begin{aligned} \min \quad & \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \sum_{i=1}^K x_i(t) c_i \\ \text{s.t.} \quad & \sum_{i=1}^K x_i(t) \mu_i \geq \lambda(t), \forall t \\ & \lim_{T \rightarrow \infty} \frac{1}{T} \sum_{t=0}^{T-1} \alpha(t) \leq MT \\ & x_i \in N, \quad \forall i, t \end{aligned}$$



Cost-Efficient Mix Scaling Algorithm

- Virtual Queue Construction $Q(t)$
- Lyapunov Drift Construction $\Delta L(t)$
- One-slot Optimization Problem Construction
- Optimization Problem Solving



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Virtual Queue

- Migration delay → **Virtual queue**
 - $Q(0) = 0$
 - $Q(t + 1) = \max\{Q(t) + \alpha(t) - MT, 0\}$
- The **equivalence** of migration delay constraint and the stability of virtual queue
 - $\lim_{T \rightarrow \infty} \sum_{t=0}^{T-1} \alpha(t) \leq MT \Leftrightarrow \lim_{T \rightarrow \infty} \frac{Q(t)}{T} = 0$
- Thus, we first construct the virtual queue and utilize it to replace the migration delay constraint



Lyapunov Drift

- To represent the stability of the virtual queue, we define two notations based on Lyapunov optimization framework
 - Lyapunov function: $L(t) = \frac{1}{2} Q(t)^2$
 - **Lyapunov drift**: $\Delta L(t) = E\{L(t+1) - L(t) | Q(t)\}$
- There always exists an **upper bound** of the Lyapunov drift:
 - $\Delta L(t) \leq M + Q(t)E\{2\frac{D(t)}{b} + B | Q(t)\}$
 - where $M = \frac{1}{2}(2\frac{D_{max}}{b} + s - MT)^2$, $B = s - MT$



One-slot Optimization Problem

- Utilizing the upper bound, we formulate the objective of the one-slot optimization problem

- $VC(t) + \Delta L(t) \leq M + VC(t) + Q(t)E\{2\frac{D(t)}{b} + B|Q(t)\}$
- where $C(t)$ is the objective of COMDT problem

- To minimize this objective, the **one-slot optimization problem** is

$$\min VC(t) + Q(t)(2\frac{D(t)}{b} + B)$$

$$s.t. \sum_{i=1}^K x_i(t)\mu_i \geq \lambda(t), \forall t$$

$$x_i \in N, \quad \forall i, t$$

- Finally, we adopt typical optimization techniques to solve it



Outline

- Problem Definition
- Cost-efficient Mix Scaling Algorithm
- Evaluation
- Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Simulation Setup

- Workload λ :
 - Generated by **TPC-W**
 - 2 types of workload: low-fluctuation & high fluctuation
- VM types: **5 types** as follows
 - capacity μ : preliminary runtime test on our OpenStack platform
 - price c : the same as AWS

Flavor	Configurations	Price/h	Price/core
m4.large	2 vCPUs, 8G RAM	\$0.979	\$0.490
m4.xlarge	4 vCPUs, 16G RAM	\$1.226	\$0.307
m4.2xlarge	8 vCPUs, 32G RAM	\$2.553	\$0.319
m4.4xlarge	16 vCPUs, 64G RAM	\$5.057	\$0.316
m4.10xlarge	40 vCPUs, 160G RAM	\$12.838	\$0.321



Comparison methods

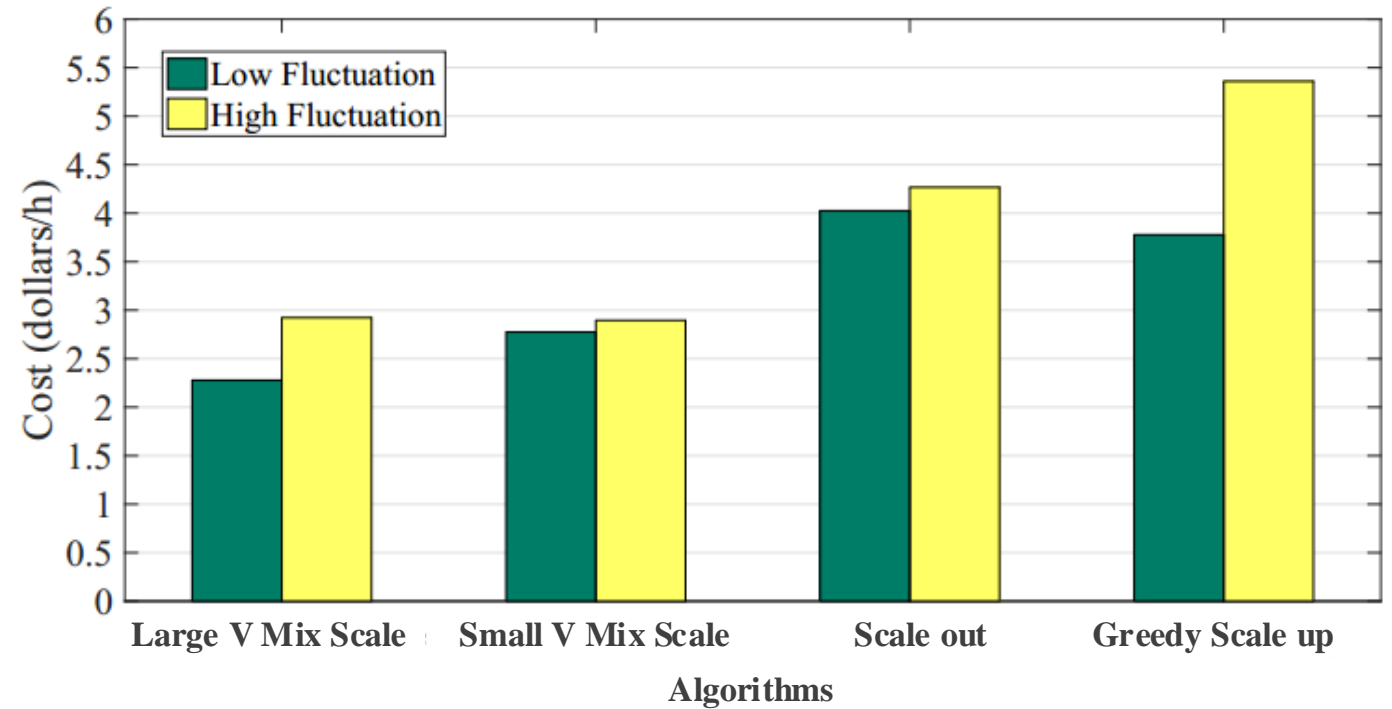
➤ 4 algorithms:

- scale out: only use **one type** VM, and scale the number of the VM
- greedy scale up: first **scale the VM type**, then the number
- mix scale: **our algorithm. 2 variations**
 - small V mix scale: focus more on migration delay
 - large V mix scale: focus more on cost




Average Cost

- Our algorithm with **small V** achieves **30.8%** and **26.3%** higher cost-efficiency than that of scale out and greedy scale up algorithms
- Our algorithm with **large V** achieves **31.1%** and **26.5%** higher cost-efficiency than that of scale out and greedy scale up algorithms




Response Time

- Under the same workload, small V mix scale algorithm can reduce **38.19%** migration delay to further reduce the response time compared with large V mix scale algorithm.



Large V



Small V



Outline

- Problem Definition
- Cost-efficient Mix Scaling Algorithm
- Evaluation
- Conclusion



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Conclusion

- Formulate the cost-migration delay tradeoff problem
 - both **cost** of cloud resources and **migration delay** are considered
- Propose the cost-efficient mix scaling algorithm
 - solve the COMDT problem utilizing the **Lyapunov optimization techniques**
- Demonstrate the efficiency and feasibility of the algorithm
 - save **31.1%** and **26.5%** cost while controlling migration delay compared with scale out and scale up algorithms



Thank you!

Q & A



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Image Size Modeling

➤ Image Size D

- $D(t) = d \sum_{x_i(t-1) > x_i(t)} |x_i(t-1) - x_i(t)|$
- where d is the average image size, and $x(t-1)$ & $x(t)$ are two constitute VM configurations



上海交通大学
SHANGHAI JIAO TONG UNIVERSITY



Tradeoff with V

- Through tuning the weight V in the one-slot optimization problem, we can control the focus between cost and migration delay:
- $V \rightarrow \infty$: reduce **more cost**, but **less control** on migration delay
 - $V \rightarrow 0$: reduce **less cost**, but **more control** on migration delay

$$\begin{aligned} \min \quad & VC(t) + Q(t)(2\frac{D(t)}{b} + B) \\ \text{s.t.} \quad & \sum_{i=1}^K x_i(t)\mu_i \geq \lambda(t), \forall t \\ & x_i \in N, \quad \forall i, t \end{aligned}$$

